

SummarAI

Individual Reflection

Smaragda

NLP Group Project · 2025–2026

Individual Reflection — Smaragda

My Contributions

Evaluation Set Design and Annotation Methodology

I designed the evaluation set the whole team used to measure the system. This included deciding on the format of `eval_set.jsonl`, writing the annotation guidelines the team followed when labeling, and coordinating labeling across all five members so everyone applied the same rules consistently.

What counts as an action item: A concrete commitment by a named or identifiable person — not a vague intention or a plan the team discussed but did not commit to.

Missing deadlines: Leave the deadline field blank rather than infer one.

Overlapping items: When one transcript has multiple overlapping action items for the same person, list each separately.

I also ensured the test set covered a realistic range of meeting types: standups, planning sessions, sales calls, retrospectives, and tricky cases (no action items at all, ambiguous ownership, hypothetical discussions). I contributed approximately 8 transcripts myself, focusing on multilingual scenarios.

Executive Summary

I wrote `deliverables/executive_summary.md` — the one-page non-technical description of the project. I led with the cost of unproductive meetings and the specific barrier multilingual teams face, then described what SummarAI does differently, then summarised the evaluation findings including the main limitation around owner attribution.

User Manual

I wrote `docs/user_manual.md` from scratch, walking through every feature: uploading a file, choosing the output language, reading the summary and takeaway bubbles, checking off action items, saving to a project, and browsing the project timeline. I also documented supported formats, known limitations, and common questions.

Slides

I put together the presentation slides (`deliverables/slides.md`), structuring the 20-minute presentation as: problem and motivation → field review → system overview → evaluation method and results → failure cases → future directions → Q&A.;

What I Learned

Building the annotation guidelines taught me that the definition of a label is itself a research decision with real consequences for the evaluation. We started with a vague shared understanding of what an action item was, and when I tried to write it down precisely, edge cases appeared immediately: is 'we should probably follow up on this' an action item? What if no one is named as responsible? What if someone agrees to do something only conditionally? These decisions determine what the model is

being evaluated against — so they determine what the evaluation actually measures. I now understand why inter-annotator agreement is reported in NLP papers: it is not a formality, it is evidence that the label definition is precise enough to be applied consistently. I also understood more deeply what the summarization task is about: not just shortening text, but identifying what is important and representing it in a structured, reusable form.