

SummarAI

Installation & Execution Guide

How to run SummarAI locally

NLP Group Project . 2025-2026

Requirements

- Python 3.10 or 3.11 (3.12+ not tested)
- pip 23+
- ffmpeg (any recent version) -- required by Whisper for audio decoding
- Anthropic API key -- needed for Claude summarisation
- At least 8 GB RAM and 3 GB free disk space for Whisper model weights

Compatible with macOS, Linux, and Windows (PowerShell).

1. Clone the Repository

```
git clone https://github.com/tono2002/nlp-project.git
cd nlp-project
```

2. Create a Virtual Environment

```
python -m venv .venv
```

macOS / Linux:

```
source .venv/bin/activate
```

Windows (PowerShell):

```
.venv\Scripts\Activate.ps1
```

3. Install Python Dependencies

```
pip install --upgrade pip
pip install -r requirements.txt
```

4. Install ffmpeg

macOS:

```
brew install ffmpeg
```

Ubuntu / Debian:

```
sudo apt install ffmpeg
```

Windows -- download from <https://ffmpeg.org/download.html> and add to PATH.

5. Configure Environment Variables

Copy the example file and add your Anthropic API key:

```
cp .env.example .env
```

Open `.env` and set:

```
ANTHROPIC_API_KEY=sk-ant-...
```

Supabase variables (`SUPABASE_URL`, `SUPABASE_ANON_KEY`) are optional. Without them, project saving is disabled but summarisation still works.

6. Run the App

```
python -m uvicorn src.app:app --reload --port 8000
```

Open **`http://127.0.0.1:8000`** in your browser.

On first run, Whisper will download the `base.en` model weights (~145 MB). This happens once and is cached for subsequent runs.

7. Optional -- Change the Whisper Model

The default model is `base.en`. For faster (less accurate) results use `tiny.en`; for better accuracy on noisy or accented audio use `small.en`:

```
WHISPER_MODEL=tiny.en python -m uvicorn src.app:app --port 8000
```

Available sizes: `tiny.en` . `base.en` . `small.en` . `medium.en`

Troubleshooting

- **ffmpeg not found:** ensure ffmpeg is on your system PATH and restart your terminal.
- **Invalid API key:** check that `ANTHROPIC_API_KEY` is set correctly in `.env`.
- **Port already in use:** kill the existing process or use a different port.
- **Slow transcription:** switch to `tiny.en` for demos; use `small.en` for better accuracy.